




46°

Lab - AWS re/Start

Uso de escalado automático en AWS (Linux)



Amazon EC2 Auto Scaling



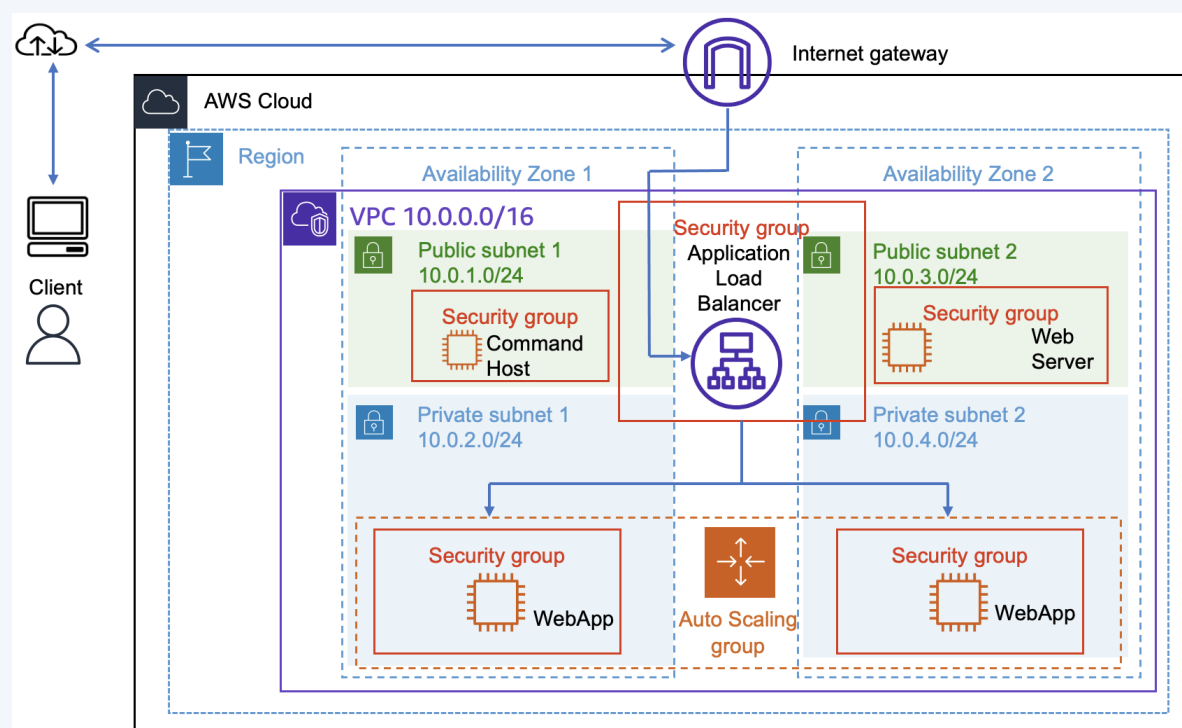
Tarea 01



Interactuando con Arquitecturas

Los objetivos son:

- Crear una instancia EC2 mediante un comando de la CLI de AWS.
- Crear una nueva AML mediante la CLI de AWS.
- Crear una plantilla de lanzamiento de Amazon EC2.
- Cree una configuración de lanzamiento de Amazon EC2 Auto Scaling.
- Configure las políticas de escalado y cree un grupo de Auto Scaling para aumentar y reducir el número de servidores en función de una carga variable.



Tarea 01



Empezamos creando la AMI de la instancia *Command Host* desde la CLI de AWS

```
us-west-2.console.aws.amazon.com/ec2-instance-connect/ssh?connType=standard&instanceId=i-086efb916c52daf02&osUser=ec2-user&region=us-west-2&sshPort=22#/  
Amazon Linux 2  
AL2 End of Life is 2025-06-30.  
A newer version of Amazon Linux is available!  
Amazon Linux 2023, GA and supported until 2028-03-15.  
https://aws.amazon.com/linux/amazon-linux-2023/  
[ec2-user@ip-10-0-1-54 ~]$ curl http://169.254.169.254/latest/dynamic/instance-identity/document | grep region  
% Total % Received % Xferd Average Speed Time Time Time Current  
Dload Upload Total Spent Left Speed  
100 474 100 474 0 0 368k 0 --:--:-- --:--:-- --:--:-- 462k  
"region": "us-west-2",  
[ec2-user@ip-10-0-1-54 ~]$ aws configure  
AWS Access Key ID [None]:  
AWS Secret Access Key [None]:  
Default region name [us-west-2]: us-west-2  
Default output format [None]: json  
[ec2-user@ip-10-0-1-54 ~]$ cd /home/ec2-user/  
[ec2-user@ip-10-0-1-54 ~]$
```

Primero, crearemos una nueva instancia, a continuación se muestra la User Data para esta instancia a crear

```
aws  
[ec2-user@ip-10-0-1-54 ~]$ more UserData.txt  
#!/bin/bash  
yum update -y --security  
amazon-linux-extras install epel -y  
yum -y install httpd php stress  
systemctl enable httpd.service  
systemctl start httpd  
cd /var/www/html  
wget http://aws-tc-largeobjects.s3.amazonaws.com/CUR-TF-100-TULABS-1/10-lab-autoscaling-linux/s3/ec2-stress.zip  
unzip ec2-stress.zip  
  
echo 'UserData has been successfully executed.' >> /home/ec2-user/result  
find -wholename /root/.history -wholename /home/*/.history -exec rm -f {} \;  
find / -name 'authorized_keys' -exec rm -f {} \;  
rm -rf /var/lib/cloud/data/scripts/*  
[ec2-user@ip-10-0-1-54 ~]$
```

Lanzamos la nueva instancia

```
us-west-2.console.aws.amazon.com/ec2-instance-connect/ssh?connType=standard&instanceId=i-086efb916c52daf02&osUser=ec2-user&region=us-west-2&sshPort=22#/  
[ec2-user@ip-10-0-1-54 ~]$ aws ec2 run-instances --key-name voykey --instance-type t3.micro --image-id ami-0d4d45428f2d4f0c --user-data file:///home/ec2-user/UserData.txt --security-group-ids sg-02ac5bfc9cca0830 --subnet-id subnet-0b92ac90c02ed95a1 --associate-public-ip-address --tag-specifications  
{  
  "Instances": [  
    {  
      "Monitoring": {  
        "State": "disabled"  
      },  
      "PublicDnsName": "",  
      "StateReason": {  
        "Message": "pending",  
        "Code": "pending"  
      },  
      "State": {  
        "Code": 0,  
        "Name": "pending"  
      },  
      "EbsOptimized": false,  
      "LaunchTime": "2024-03-05T22:48:40.000Z",  
      "PrivateIpAddress": "10.0.3.206",  
      "ProductCodes": [],  
      "OwnerId": "ac-07c536590dcad87be",  
      "CpuOptions": {  
        "CoreCount": 1,  
        "ThreadsPerCore": 2  
      },  
      "StateTransitionReason": "",  
      "InstanceId": "i-07f2e34d16ddcd5d59",  
      "EnableSupport": true,  
      "ImageId": "ami-0d4d45428f2d4f0c",  
      "PrivateDnsName": "ip-10-0-3-206.us-west-2.compute.internal",  
      "KeyName": "voykey",  
      "SecurityGroups": [  
        {  
          "GroupName": "HTTPAccess",  
          "GroupId": "sg-02ac5bfc9cca08783"  
        }  
      ],  
      "ClientToken": "8dd2f813-c6fa-407e-bb9c-fb1cf1908906",  
      "SubnetId": "subnet-0b92ac90c02ed95a1"  
    }  
  ]  
}
```


Tarea 01



Ahora, veamos esta nueva instancia:



Luego, procedemos a crear una AMI de esta instancia



Después de ello, procedemos a crear un ELB:

Create Application Load Balancer [Info](#)

The Application Load Balancer distributes incoming HTTP and HTTPS traffic across multiple targets such as Amazon EC2 instances, microservices, and containers, based on request attributes. When the load balancer receives a connection request, it evaluates the listener rules in priority order to determine which rule to apply, and if applicable, it selects a target from the target group for the rule action.

► How Application Load Balancers work

Basic configuration

Load balancer name

Name must be unique within your AWS account and can't be changed after the load balancer is created.

WebServerELB

A maximum of 32 alphanumeric characters including hyphens are allowed, but the name must not begin or end with a hyphen.

Scheme [Info](#)

Scheme can't be changed after the load balancer is created.

☒ Internet-facing

An internet-facing load balancer routes requests from clients over the internet to targets. Requires a public subnet. [Learn more](#)

☐ Internal

An internal load balancer routes requests from clients to targets using private IP addresses.

IP address type [Info](#)

Select the type of IP addresses that your subnets use.

☒ IPv4

Recommended for internal load balancers.

☐ Dualstack

Includes IPv4 and IPv6 addresses.

Network mapping [Info](#)

The load balancer routes traffic to targets in the selected subnets, and in accordance with your IP address settings.

VPC [Info](#)

Select the virtual private cloud (VPC) for your targets or you can [create a new VPC](#). Only VPCs with an internet gateway are enabled for selection. The selected VPC can't be changed after the load balancer is created. To confirm the VPC for your targets, view your [target groups](#).

Lab VPC

vpc-07c536590dcad87be

IPv4: 10.0.0.0/16

Mappings [Info](#)

Select at least two Availability Zones and one subnet per zone. The load balancer routes traffic to targets in these Availability Zones only. Availability Zones that are not supported by the load balancer or the VPC are not available for selection.

☒ us-west-2a (usw2-az2)

Subnet

subnet-06f2def419b93379c

Public Subnet 1

IPv4 address

Assigned by AWS

☒ us-west-2b (usw2-az1)

Subnet

subnet-0b92ac90c02ed95a1

Public Subnet 2

IPv4 address

Assigned by AWS

Tarea 01



Más de la configuración del ALB (Application Load Balancer)

EC2 > Target groups > Create target group

Step 1
Specify group details

Step 2
Register targets

Specify group details

Your load balancer routes requests to the targets in a target group and performs health checks on the targets.

Basic configuration

Settings in this section can't be changed after the target group is created.

Choose a target type

Instances

- Supports load balancing to instances within a specific VPC.
- Facilitates the use of [Amazon EC2 Auto Scaling](#) to manage and scale your EC2 capacity.

IP addresses

- Supports load balancing to VPC and on-premises resources.
- Facilitates routing to multiple IP addresses and network interfaces on the same instance.
- Offers flexibility with microservice based architectures, simplifying inter-application communication.
- Supports IPv6 targets, enabling end-to-end IPv6 communication, and IPv4-to-IPv6 NAT.

Lambda function

- Facilitates routing to a single Lambda function.
- Accessible to Application Load Balancers only.

Application Load Balancer

- Offers the flexibility for a Network Load Balancer to accept and route TCP requests within a specific VPC.
- Facilitates using static IP addresses and PrivateLink with an Application Load Balancer.

Target group name

webserver-app

Listeners and routing

A listener is a process that checks for connection requests using the port and protocol you configure. The rules that you define for a listener determine how the load balancer routes requests to its registered targets.

▼ Listener HTTP:80

Remove

Protocol

Port

Default action

Info

HTTP

:

80

1-65535

Forward to

webserver-app

HTTP

↺

Create target group

Listener tags - optional

Consider adding tags to your listener. Tags enable you to categorize your AWS resources so you can more easily manage them.

Add listener tag

You can add up to 50 more tags.

Add listener

Health checks

The associated load balancer periodically sends requests, per the settings below, to the registered targets to test their status.

Health check protocol

HTTP

Health check path

Use the default path of "/" to perform health checks on the root, or specify a custom path if preferred.

/index.php

Up to 1024 characters allowed.

► Advanced health check settings

Luego, procedemos a crear la plantilla de lanzamiento

▼ Summary

Software Image (AMI)

WebServerAMI

ami-07b04e1cfb96876cd

Virtual server type (instance type)

t3.micro

Firewall (security group)

HTTPAccess

Storage (volumes)

1 volume(s) - 8 GiB

Cancel

Create launch template

SWIPE

Tarea 01



Luego creamos el ASG (Auto Scaling Group)

Instance type requirements [Info](#)

Override launch template

You can keep the same instance attributes or instance type from your launch template, or you can choose to override the launch template by specifying different instance attributes or manually adding instance types.

Launch template	Version	Description
web-app-launch-template ↗	Default	web-app-launch-template
lt-01bc43cbcc2d4a863		
Instance type		
t3.micro		

Network [Info](#)

For most applications, you can use multiple Availability Zones and let EC2 Auto Scaling balance your instances across the zones. The default VPC and default subnets are suitable for getting started quickly.

VPC

Choose the VPC that defines the virtual network for your Auto Scaling group.

vpc-07c536590dcad87be (Lab VPC)

10.0.0.0/16

↻

[Create a VPC](#) [↗](#)

Availability Zones and subnets

Define which Availability Zones and subnets your Auto Scaling group can use in the chosen VPC.

Select Availability Zones and subnets

↻

↻

us-west-2a | subnet-0cc436d2de1674e74 (Private Subnet 1)

10.0.2.0/24

✕

us-west-2b | subnet-00dfef2f5052e3e74 (Private Subnet 2)

10.0.4.0/24

✕

[Create a subnet](#) [↗](#)

Health checks

Health checks increase availability by replacing unhealthy instances. When you use multiple health checks, all are evaluated, and if at least one fails, instance replacement occurs.

EC2 health checks

Always enabled

↻

Additional health check types - *optional* | [Info](#)

Turn on Elastic Load Balancing health checks

Recommended

Elastic Load Balancing monitors whether instances are available to handle requests. When it reports an unhealthy instance, EC2 Auto Scaling can replace it on its next periodic check.

EC2 Auto Scaling will start to detect and act on health checks performed by Elastic Load Balancing.

✕

To avoid unexpected terminations, first verify the settings of these health checks in the [Load Balancer console](#) [↗](#)

Turn on VPC Lattice health checks

VPC Lattice can monitor whether instances are available to handle requests. If it considers a target as failed a health check, EC2 Auto Scaling replaces it after its next periodic check.

Health check grace period | [Info](#)

This time period delays the first health check until your instances finish initializing. It doesn't prevent an instance from terminating when placed into a non-running state.

300

seconds

Load balancing [Info](#)

Use the options below to attach your Auto Scaling group to an existing load balancer, or to a new load balancer that you define.

☐ No load balancer

Traffic to your Auto Scaling group will not be fronted by a load balancer.

☒ Attach to an existing load balancer

Choose from your existing load balancers.

☐ Attach to a new load balancer

Quickly create a basic load balancer to attach to your Auto Scaling group.

Attach to an existing load balancer

Select the load balancers that you want to attach to your Auto Scaling group.

☒ Choose from your load balancer target groups

This option allows you to attach Application, Network, or Gateway Load Balancers.

☐ Choose from Classic Load Balancers

Existing load balancer target groups

Only instance target groups that belong to the same VPC as your Auto Scaling group are available for selection.

Select target groups

↻

↻

webserver-app | HTTP

✕

Application Load Balancer: WebServerELB

Scaling [Info](#)

You can resize your Auto Scaling group manually or automatically to meet changes in demand.

Scaling limits

Set limits on how much your desired capacity can be increased or decreased.

Min desired capacity

2

Equal or less than desired capacity

Max desired capacity

4

Equal or greater than desired capacity

Automatic scaling - *optional*

Choose whether to use a target tracking policy | [Info](#)

You can set up other metric-based scaling policies and scheduled scaling after creating your Auto Scaling group.

☐ No scaling policies

Your Auto Scaling group will remain at its initial size and will not dynamically resize to meet demand.

☒ Target tracking scaling policy

Choose a CloudWatch metric and target value and let the scaling policy adjust the desired capacity in proportion to the metric's value.

Scaling policy name

Target Tracking Policy

Metric type | [Info](#)

Monitored metric that determines if resource utilization is too low or high. If using EC2 metrics, consider enabling detailed monitoring for better scaling performance.

Average CPU utilization

▼

Target value

50

Instance warmup | [Info](#)

300


seconds

Tarea 01



Y lo probamos, funciona correctamente

Generating load



Generate Load

Start Stress

Stop Stress

CloudWatch > Alarms

Alarms (2)

☐ Hide Auto Scaling alarms

Clear selection

Create composite alarm

Actions

Create alarm

Search

Alarm state: Any

Alarm type: Any

Actions status: Any

< 1 >

<input type="checkbox"/>	Name	State	Last state update	Conditions	Actions
<input type="checkbox"/>	TargetTracking-Web App Auto Scaling Group-AlarmLow-b9c62411-bb60-4813-917b-6ea3507efbc0	In alarm	2024-03-06 00:36:03	CPUUtilization < 35 for 15 datapoints within 15 minutes	Actions enabled
<input type="checkbox"/>	TargetTracking-Web App Auto Scaling Group-AlarmHigh-e5c56949-df71-42a1-adab-a4ea50c6e330	OK	2024-03-06 00:23:06	CPUUtilization > 50 for 3 datapoints within 3 minutes	Actions enabled

Activity history (4)

Filter activity history

< 1 >

Status	Description	Cause	Start time	End time
Successful	Terminating EC2 instance: i-0a916fd1a4ebebba0	At 2024-03-06T00:36:04Z a monitor alarm TargetTracking-Web App Auto Scaling Group-AlarmLow-b9c62411-bb60-4813-917b-6ea3507efbc0 in state ALARM triggered policy Target Tracking Policy changing the desired capacity from 3 to 2. At 2024-03-06T00:36:15Z an instance was taken out of service in response to a difference between desired and actual capacity, shrinking the capacity from 3 to 2. At 2024-03-06T00:36:15Z instance i-0a916fd1a4ebebba0 was selected for termination.	2024 March 05, 07:36:15 PM -05:00	2024 March 05, 07:42:22 PM -05:00
Successful	Launching a new EC2 instance: i-068cbc849f9ff492e	At 2024-03-06T00:19:06Z a monitor alarm TargetTracking-Web App Auto Scaling Group-AlarmHigh-e5c56949-df71-42a1-adab-a4ea50c6e330 in state ALARM triggered policy Target Tracking Policy changing the desired capacity from 2 to 3. At 2024-03-06T00:19:09Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 2 to 3.	2024 March 05, 07:19:11 PM -05:00	2024 March 05, 07:24:17 PM -05:00
Successful	Launching a new EC2 instance: i-03d758b2c2f928ef3	At 2024-03-06T00:01:48Z a user request created an AutoScalingGroup changing the desired capacity from 0 to 2. At 2024-03-06T00:01:53Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 0 to 2.	2024 March 05, 07:01:55 PM -05:00	2024 March 05, 07:02:01 PM -05:00
Successful	Launching a new EC2 instance: i-0a916fd1a4ebebba0	At 2024-03-06T00:01:48Z a user request created an AutoScalingGroup changing the desired capacity from 0 to 2. At 2024-03-06T00:01:53Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 0 to 2.	2024 March 05, 07:01:55 PM -05:00	2024 March 05, 07:02:01 PM -05:00